

```
In [4]: import pandas as pd
```

1) Read the dataset

```
In [5]: df=pd.read_csv("Churn_Modelling.csv")  
# 0 means not exit and 1 means exit
```

```
In [6]: df.head(4)
```

```
Out [6]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0

```
In [7]: df.shape
```

```
Out [7]: (10000, 14)
```

2) Distinguish the feature and target set and divide the data set into training and test sets.

```
In [8]: df.columns
```

```
Out [8]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',  
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',  
              'IsActiveMember', 'EstimatedSalary', 'Exited'],  
              dtype='object')
```

```
In [9]: X=df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'Esti  
Y=df['Exited']
```

```
In [10]: Y.value_counts()
```

```
Out [10]: Exited  
0      7963  
1      2037  
Name: count, dtype: int64
```

3) Normalize the train and test data.

```
In [17]: from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()  
X_scaled=scaler.fit_transform(X)
```

```
In [18]: X_scaled
```

```
Out [18]: array([[ -0.32622142,  0.29351742, -1.04175968, ...,  0.64609167,  
                  0.97024255,  0.02188649],  
                [ -0.44003595,  0.19816383, -1.38753759, ..., -1.54776799,  
                  0.97024255,  0.21653375],  
                [ -1.53679418,  0.29351742,  1.03290776, ...,  0.64609167,  
                 -1.03067011,  0.2406869 ],  
                ...,  
                [  0.60498839, -0.27860412,  0.68712986, ..., -1.54776799,  
                  0.97024255, -1.00864308],  
                [  1.25683526,  0.29351742, -0.69598177, ...,  0.64609167,  
                 -1.03067011, -0.12523071],  
                [  1.46377078, -1.04143285, -0.35020386, ...,  0.64609167,  
                 -1.03067011, -1.07636976]])
```

4) Initialize and build the model. Identify the points of improvement and implement the same.

```
In [19]: from sklearn.model_selection import train_test_split  
X_train,X_test,Y_train,Y_test=train_test_split(X_scaled,Y,test_size=0.25,random_state=0)
```

```
In [20]: from sklearn.neural_network import MLPClassifier  
nn_model=MLPClassifier(hidden_layer_sizes=(80,80,80),random_state=0,max_iter=100,activation='relu')
```

```
In [21]: # MLPClassifier stands for Multilayer Perceptron Classifier. It is a type of neural network that consists of an input layer, one or more hidden layers, and an output layer.
# The tuple (80, 80, 80) indicates that the neural network will have three hidden layers, each containing 80 neurons.
# max_iter sets the maximum number of iterations (epochs) the optimizer will perform during training.
# activation specifies the activation function to use for the hidden layers. 'relu': Rectified Linear Unit (ReLU) activation function.
```

```
In [22]: nn_model.fit(X_train,Y_train)
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\normalization\_maxabs.py:686:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.
warnings.warn(
```

```
Out [22]: MLPClassifier
MLPClassifier(hidden_layer_sizes=(80, 80, 80), max_iter=100, random_state=0)
```

```
In [23]: y_pred=nn_model.predict(X_test)
y_pred
```

```
Out [23]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```


```
In [24]: from sklearn.metrics import ConfusionMatrixDisplay, classification_report
from sklearn.metrics import accuracy_score
```

```
In [25]: Y_test.value_counts()
```

```
Out [25]: Exited
0      1991
1       509
Name: count, dtype: int64
```

```
In [26]: ConfusionMatrixDisplay.from_predictions(Y_test,y_pred)
```

```
Out [26]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1f9f58d7550>
```



```
In [27]: accuracy_score(Y_test,y_pred)
```

```
Out [27]: 0.8316
```

```
In [28]: print(classification_report(Y_test,y_pred))
```

```

              precision    recall  f1-score   support

     0       0.88        0.91        0.90        1991
     1       0.60        0.51        0.55         509

 accuracy: 0.83
macro avg: 0.74        0.71        0.72        2500
weighted avg: 0.82        0.83        0.83        2500
```

```
In [29]: # the recall value for 0 (not exit) is 0.81 whereas recall value for 1 (exit) is 0.24, so there is high imbalance
# the gap should be as less as possible.
# this happens because of imbalance dataset, we can see that entries for 0 are very more compared to 1
# So, we need to balance the dataset
```

```
In [30]: Y.value_counts()
```

```
Out [30]: Exited  
0      7963  
1      2037  
Name: count, dtype: int64
```

```
In [31]: # here we need to increase entries of 1
```

```
In [3]: pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in c:\users\ashvini mahajan\anaconda\lib\site-packages (0.11.0)  
Requirement already satisfied: numpy>=1.17.3 in c:\users\ashvini mahajan\anaconda\lib\site-packages (from imbalanced-learn) (1.26.4)  
Requirement already satisfied: scipy>=1.5.0 in c:\users\ashvini mahajan\anaconda\lib\site-packages (from imbalanced-learn) (1.11.4)  
Requirement already satisfied: scikit-learn>=1.0.2 in c:\users\ashvini mahajan\anaconda\lib\site-packages (from imbalanced-learn) (1.2.2)  
Requirement already satisfied: joblib>=1.1.1 in c:\users\ashvini mahajan\anaconda\lib\site-packages (from imbalanced-learn) (1.1.1)  
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\ashvini mahajan\anaconda\lib\site-packages (from imbalanced-learn) (2.2.1)  
Note: you may need to restart the kernel to use updated packages.
```

```
In [14]: from imblearn.over_sampling import RandomOverSampler  
ros=RandomOverSampler(random_state=0)  
X_new,Y_new=ros.fit_resample(X,Y)
```

```
In [15]: Y_new.value_counts()
```

```
Out [15]: Exited  
1      7963  
0      7963  
Name: count, dtype: int64
```

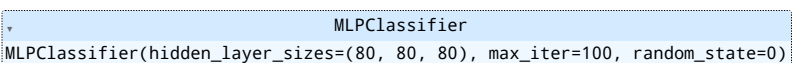
```
In [16]: # now dataset is balanced  
# now again make the model
```

```
In [32]: from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()  
X_scaled=scaler.fit_transform(X_new)
```

```
In [33]: Xn_train,Xn_test,Yn_train,Yn_test =train_test_split(X_scaled,Y_new,random_state=0,test_size=0.25)
```

```
In [34]: from sklearn.neural_network import MLPClassifier  
nn_model=MLPClassifier(hidden_layer_sizes=(80,80,80),random_state=0,max_iter=100,activation='relu')  
nn_model.fit(Xn_train,Yn_train)
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:686:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.  
warnings.warn(
```

```
Out [34]: MLPClassifier  
MLPClassifier(hidden_layer_sizes=(80, 80, 80), max_iter=100, random_state=0)
```

```
In [35]: y_pred=nn_model.predict(Xn_test)  
y_pred
```


```
Out [35]: array([1, 1, 1, ..., 1, 1, 0], dtype=int64)
```

```
In [37]: Yn_test.value_counts()
```

```
Out [37]: Exited  
1      2001  
0      1981  
Name: count, dtype: int64
```

```
In [38]: ConfusionMatrixDisplay.from_predictions(Yn_test,y_pred)
```

```
Out [38]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1f9f73d3d50>
```



```
In [40]: accuracy_score(Yn_test,y_pred)
```

```
Out [40]: 0.8455549974886991
```

```
In [42]: print(classification_report(Yn_test,y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.78	0.83	1981
1	0.81	0.91	0.86	2001
accuracy			0.85	3982
macro avg	0.85	0.85	0.84	3982
weighted avg	0.85	0.85	0.84	3982

```
In [ ]:
```